

THE AESTHETICS OF AI-WRITTEN TEXTS

HOW THIS ANTHOLOGY EVOLVED

My interest in machine-generated texts goes back a long way. As a first-year student, I read Max Bense's seminal book *Einführung in die informationstheoretische Ästhetik: Grundlegung und Anwendung in der Texttheorie* (1971). Bense cites, among others, "Stochastische Texte" by Theo Lutz. They made a huge impression on me. And, as soon as I had my first computer, I tried to emulate this method of text production.

In the 1990s, I read books about LISP and learned programming in PROLOG. Both languages were then considered essential for studying Artificial Intelligence. My biggest success in PROLOG was a program that could derive the name of the metre from the sequence of syllables of a verse. It was then, in the early 1990s, that I gave a short talk on rule-based text generation that ended with the prediction: one day computers will be able to write whole stories without human intervention.

That day has arrived. In early 2022, a friend asked me if it would be possible to generate a blog post with AI. I tried a few text-completion algorithms, but the results were too weird for serious blogging. Then, in November 2022, OpenAI released GPT-3.5 (ChatGPT), and the world of machine writing changed in a fundamental way: for the first time it was possible to generate texts that were more than experiments for nerds.

ChatGPT triggered a worldwide hype. Suddenly everybody talked about AI. And nearly everybody used it. But, from day one, there was a sharp divide: people who welcomed ChatGPT as the forerunner of an AGI or even ASI that would make life easier for everybody, if not bring paradise on Earth, and people who saw in the chatbot a harbinger of doom. The divide persists to the present day.

Criticism of AI has many important points to make:

- Training Large Language Models requires enormous amounts of electrical energy.
- The data used in training current models were collected without asking copyright owners for permission.
- Only a few states (the USA, China, the UAE) and a very limited number of companies (Google, Amazon, Meta, xAI, OpenAI, Anthropic, Moonshot AI, High-Flyer) have the money and know-how needed to play in the first league of AI research and development.
- Large Language Models reproduce biases common on the internet.
- English is by far the dominant language on the internet and thus is overrepresented in AI training data.
- AI is a threat to millions of jobs.
- AI is used to flood the internet with fake news, and AI bots are the nemesis of social media.
- The military use of AI will have disastrous consequences.
- AGI and then ASI might develop into an existential threat for humanity.

- The use of AI tools makes people lazy and, over time, leads to a deterioration of their mental capacities.
- If children or students are granted access to AI tools, this will hinder their mental and intellectual development.
- People use AI for guidance, but AI doesn't know right from wrong or truth from falsehood.
- AI is not a person, but there are people who enter into an emotional relationship with a chatbot.

None of these objections to AI is without merit, and every single one of them needs to be addressed if AI is to be beneficial for all of humanity.

There is one argument, however, that is often brought forward and which I think is without merit:

- AI is a bad writer.

On the contrary: AI is an excellent writer. In fact, it writes better than most of us. From its enormous corpus of training data, it has learned many of the rules, finer points, and intricacies of human language—not only one language but many. It is capable of writing in English, German, French, Spanish, Turkish, Arabic, Urdu, Hindi, Japanese, Mandarin, Korean, Vietnamese, and so on and so forth. True: AI is not particularly good at Mohawk, Navajo, Quechua, Uzbek, Kazakh, Tamazight, or many languages from south of the Sahara. But in many cases, as with te reo Māori, the Indigenous language of Aotearoa (New Zealand), efforts are being made to close the gap, and the day is not far when AI will be able to understand and write/speak in any language a user might want to converse in.

With an AI chatbot you can even chat in Latin, Ancient Greek, Ottoman Turkish or—maybe at some point in the future—Aztec or Akkadian.

So the situation we face today (early 2026) is that humans have a rival in the arena of language. Not many people are fully aware that this is the case, but a great many can feel it. For me, this explains, to a certain degree, why the claim “AI is a bad writer” is so popular.

What has all this got to do with our anthology of machine-written texts?

This anthology is not meant to accuse or defend. It presents texts generated by Large Language Models “as is,” without human intervention other than the prompt that initiates the generation of a text. The output will be presented raw and unedited. It is up to the reader to find errors, inconsistencies, strange artifacts (like, for instance, Chinese ideograms in an English sentence), overused em dashes, and more.

But first and foremost, it is meant to explore the current state of machine writing. Let us inspect these texts that have no author, no intent, no message, no soul but still succeed in eliciting in the reader what Roland Barthes calls “le plaisir du texte.”

Constraints

Machine-generated texts are swamping the internet and social media. As a rule, they tend to conceal their origins and masquerade as products of genuine human effort. The sheer number of these texts, and the difficulty of identifying them, make it impossible to obtain a representative sample. But we are not statisticians; we are readers who enjoy well-written prose.

This anthology therefore presents texts that belong to the broad category of fiction. To focus on factors such as fluency, style, grammatical correctness, and other linguistic features, we limited the corpus to texts in English and excluded both texts in other languages and translations. A second anthology featuring texts in languages beyond English is in the early stages of planning.

Author: mtg

Copyright: [see copyright notice](#)

[Back to TOC](#)

Preface

This anthology presents a selection of texts generated by state-of-the-art large language models (LLMs). It is not a celebration of artificial intelligence, nor is it an attempt to demonstrate technological prowess. Rather, it is a curated collection designed to establish and model standards for how sophisticated machine-generated text should be treated, contextualized, and presented within the literary and intellectual landscape.

Historical Context and Purpose

The history of machine-written text spans nearly eight decades, from Christopher Strachey's algorithmic love letters in 1952 to contemporary transformer-based language models. Yet the proliferation of unvetted, low-quality AI-generated content in recent years has obscured an important distinction: not all machine-generated text merits publication. Only work of genuine sophistication—texts that demonstrate coherence, originality, and aesthetic or intellectual value—warrants serious engagement.

Fundamental Principles

This anthology operates under three fundamental principles:

1. **Curatorial Rigor:** Every text included has been carefully selected and evaluated against rigorous standards. We prioritize selections that successfully mirror the cadences and nuances of human writing, focusing on prose that transcends simple data output to achieve a recognizable literary voice. The presence of human judgment and editorial contribution is not incidental; it is essential.
2. **Transparency About Authorship:** Each text is labeled according to its provenance—AI-generated only, human-written only, or human-assisted by AI. This labeling extends to legal status: AI-only content is released as public domain; human-written material is copyright protected; mixed authorship occupies a legal grey zone that reflects the unsettled questions surrounding intellectual property in the age of AI.
3. **Contextual Framing:** Human-written introductions, notes, and commentary accompany the machine-generated texts. This paratext is not mere scaffolding; it constitutes the curatorial layer through which readers encounter and interpret the work. By presenting a wide array of genres, we allow readers to observe how effectively these models navigate the distinct conventions of different forms, from lyrical essays to technical reports.

The Human Core

The legal framework deserves particular attention. In most jurisdictions, AI-generated text alone cannot be copyrighted. However, substantial human curation, selection, and commentary can secure copyright protection for an anthology as a whole. This structure is not a legal technicality but an intellectual

statement: it acknowledges that human judgment—in choosing what deserves publication, how to present it, and what it means—remains irreplaceable.

Conclusion

The texts that follow invite readers to consider what machines can now accomplish and, equally important, what they cannot. They are offered not as curiosities but as serious contributions to contemporary literature and thought.

[Back to TOC](#)

Introduction

Synthetic Texts

An Anthology

Curated by Michael Glünz

Published by Calya Digital Publishing Services
an imprint of Calya Consult GmbH, Zug, Schweiz

Introduction

Can AI write?

Writing as Text Production

In the technical realm, a text is a sequence of characters, encoded as ASCII, UTF-8, UTF-16, UTF-32, or any other form of character encoding. A character is defined as element of a character set, the non-technical term being "alphabet." Character sets are limited, the Latin Alphabet, for instance, has 26 characters (letters), the Arabic alphabet has two or three more. As of September 2025, the largest character set in use is called Unicode 17.0.0. It defines a set of 159,801 characters, among which 101'996, nearly two-thirds, are CJK (Chinese, Japanese, and Korean) ideographs. But even when an alphabet has a size of one hundred thousand it is still limited and small in comparison to the size of the texts that can be generated based on this, or any other, alphabet.

A text is a sequence of characters, and every character represents a choice made among the letters of the alphabet. If my text starts with 'A' this means I selected 'A' from {'A', 'B', 'C' ... 'X', 'Y', 'Z'}. For the next letter of my text, I repeat the selection process, choosing, for instance, 'N' and then 'T' and so on. Every position in my text has 26—for the sake of simplicity, we stay with the Classical Latin all-uppercase alphabet—items to choose from and thus, for two positions we have 26 times 26 possible combinations, for three positions 26 times 26 times 26 (26 to the power of 3) combinations which equals 17,576 possible texts of three letters length.

The number of combinations grows very fast and a text of 40 lines with 80 characters each has such a vast number of possible combinations of letters that our universe would be much too small to hold all the pages we could fill with them.

If we let a machine choose characters at random we'll most likely end up with something like this:

DMBTBQQEOBNCUKJLMNFFLWDWHGQOINRLLEBJHGFYXKPVKBZEHEXIVLM0JEBHRRDFMUEGDWZTEKBOHC
LNLDA XVLEUUURCXEZASCMKPUBLAGZQONXUBOLEVDFKXUGDSSMUHSZQKIFEGCCXDEJPPMDWIIIBRYSWTDV
BXNJKVZWREDXZRCKTJJZVAQDBUXYEHFLZKUPNTDUWATGLEHOTNWRVOQEKMGCISSVIDBUHLCBIPDIIWO

Technically, every sequence of characters counts as 'text' and can be saved on a computer disk as '.txt' file. The definition of text per se does not restrict the length of a text but in the real world we inhabit the constraints of storage and processing put strict limits on how long our texts can become.

In order to handle text, people who wrote them had to break down the sequence of letters into lines, pages, and books. So we can, for example, have a book with 410 pages, each of which has 40 lines at 80 characters, bringing the total to 1,312,000 characters.

The Library of Babel

In 1941, the Argentine author and librarian Jorge Luis Borges (1899–1986) published a short story under the title: *La biblioteca de Babel*. In this story, he imagines a vast library filled with books of uniform size, the 410 pages we just mentioned, and convincingly argues that somewhere in this library there must be a book that contains your biography from birth to death with all the exact dates and places. But there will also be false biographies of you or your biography but with another name attached to it and so on.

The problem we, as readers, have with this library is that we can be sure a certain text is there but we have no way to find it. It would take an amount of time greater than the lifespan of the universe to scan the volumes and pages, and we still couldn't be sure where to find our biography or tomorrow's sports results. We'd be facing an ocean of gibberish that conceals the one pearl we are looking for.

Yes, the Library of Babel is there, at least as an idea, but it is a forbidding place, an endless labyrinth where we get lost forever, searching for meaning and finding nothing but random noise.

An Army of Typing Monkeys

The idea of a Library of Babel is closely related to the "infinite monkey theorem" which says that a monkey—or a whole army of monkeys—hitting keys independently and at random on a typewriter keyboard for an infinite amount of time will almost surely type any given text, including the complete works of William Shakespeare or your own biography. However, a study published in 2024 by Stephen Woodcock and Jay Falletta took a closer look at the numbers and reached the conclusion that in a finite universe there is simply not enough time, nor are there enough monkeys, to generate even a relatively short non-trivial text by letting the monkeys do the typing.

Reference: [A numerical evaluation of the Finite Monkeys Theorem](#)

To speed up text production, we nowadays have computers that are able to generate vast amounts of text almost in an instant but even with the fastest computers we would still not come anywhere near the immense number of possible combinations of letters on a page, let alone the number of books filled with such pages.

And so, since finding islands of meaning in an ocean of random gibberish is practically impossible, we have to look elsewhere for a way to navigate the rough seas.

Machines that write

From what we have discussed so far, it should be clear that a definition of text as sequence of characters is far too wide to be useful in the context of writing. Writing is more than randomly stringing characters together. Writing produces texts that have form and structure, obey rules, and finally yield meaning.

Programmable computers first appeared in the 1940s with Konrad Zuse's Z3 (1941) and Tommy Flowers's Colossus Mark 1 (1943). A decade later, in 1952, Christopher Strachey wrote a computer program that generated love letters, the first of its kind. In 1959, the German computer scientist Theo Lutz used a Zuse Z22 computer to write "Stochastische Texte," sequences of words that resembled poems.

The year 1966 marks the creation of two other memorable projects of computer generated text:

- Joseph Weizenbaum's *ELIZA*: a clever interactive algorithm that simulates a dialogue with its user; to Weizenbaum's surprise, some users of *ELIZA* entered into discussions of real-life problems with this algorithmic "therapist."
- Nanni Balestrini's *Tristano*: an experimental generative novel that was conceived to be read by each reader differently, since each section is randomly shuffled. Owing to the technical limitations of 1960s printing presses, it took several decades for the author's vision to materialize as a series of printed books, each featuring a unique arrangement of sections.

Reference:

Lutz [Stochastische Texte](#)

Balestrini [Tristano](#)

Ever since John McCarthy coined the term "artificial intelligence" (AI) in 1956, researchers in this field have tried to create algorithms that are able to understand human language and produce output that looks like human language and can be understood by human readers.

The early attempts at text generation resulted in character sequences that had, undeniably, much in common with text written by humans: the sequences were divided into words and sentences, the words were taken from existing dictionaries, the sentences followed the rules of grammar. But, in cases where they were the result of stochastic processes, the texts didn't make much sense or, in cases where fixed templates were used, they were all variations of one and the same structure.

Noam Chomsky's 1957 book "Syntactic Structures" marks a breakthrough in structural linguistics. Chomsky separates language as an abstract system (competence) from the concrete utterances of human speakers (performance) and defines syntax as a finite set of rules capable of generating an unlimited number of sentences. Chomsky's work not only had a profound influence on theories of human language but it also laid the foundations for the development of programming languages, i.e., formal systems that transform human readable program input into the executable machine code of computers.

For some time, culminating in the 1980s, rules were thought to be the solution for understanding language as well as solving many other kinds of problem. A lot of research went into "expert systems." An expert system is a collection of rules that, when applied in the right sequence to a basis of known facts, can map a question to its correct answer. In 1972, Alain Colmerauer and Philippe Roussel at the University of Aix-Marseille in France developed a special programming language for the PROcessing of LOGical problems and named it PROLOG. In numerous cases, it proved its eminent value as a tool for implementing expert systems but it missed its original goal, namely to facilitate language processing.

In the 1990s, the enthusiasm that had accompanied visions of artificial intelligence based on expert systems and rule-based problem solving cooled off and most research activity went into a different direction. Human language in all its complexity seemed out of reach for computers. Speech synthesis resulted in unnatural sounding voices, speech recognition reached its limits when more than one speaker or more than simple sentences were involved, and machine translation produced awkward, sometimes even laughable results.

But slowly, new methods of studying language gained traction. Researchers in the field of Natural Language Processing (NLP) experimented with neural networks and statistical approaches like Markov models or Bayesian networks. And with the exponential growth of storage capacity, computer memory, and processing power, new approaches to some of the hitherto insurmountable barriers were, at last, successful. Statistics and neural networks proved to be the keys to the kingdom. They required massive amounts of memory space and huge numbers of calculations but this brute force attack paid out. Finally, with the introduction of the transformer architecture, the machines learned to speak and understand human language.

In today's parlance, outside the academic circles of AI researchers, the term "Artificial Intelligence" (AI) is used more or less as a synonym for "Large Language Model" (LLM). An LLM is a statistical model of how different aspects of language are related to each other. They are usually based on very large amounts of text, going into the billions of individual documents. In a process called "training," the documents are broken down into pieces called "tokens"-which can be words or parts of words-and represented by mathematical vectors. These vectors are fed into gigantic neural networks with billions of "weights" (also called "parameters") which are then repeatedly adjusted until the accuracy of predicting the next token in a sequence of tokens satisfies the expected value.

At their core, Large Language Models are prediction machines. For a given sequence of tokens, they calculate the probability distribution for the next token based on patterns found in their training data. Ideally, this calculation would be deterministic and behave like a mathematical function, i.e., the same

input will lead to the same output. But in practice, small variations caused by the hardware and software used in the vast number of calculations lead to much bigger variations in the calculated output. In addition to this kind of uncontrolled variations LLMs can also introduce controlled randomness. The degree of such randomness is called "temperature" and its aim is to make the output more interesting.

LLMs are built to generate verbal output as response to verbal input. To this end, they draw on a vast reservoir of human texts. They don't follow linguistic rules but calculate the most probable next token. Due to the size of the data and to the multidimensionality of their vectors, the outcome is more often than not a response that shows all features of a well-formed text. Without formulating any rules, the user can via verbal prompt impose constraints on the response. Such constraint can be anything like genre, style, voice, tone, sentiment, length of text, and so on and so forth. The rules and conventions of genre and style, for instance, don't need to be stated explicitly; they are implicit in the data and their application is simulated by calculating probabilities.

Their ease of use and their enormous versatility made LLMs an instant success with the general public. You don't have to be an expert, you don't even have to be able to write complete sentences like an educated person. All you have to do is type something the LLM, through the interface of a chatbot, can analyze and interpret as a prompt. You could enter the words: "write funny limerick easter bunny" and the chatbot would understand this as a request for a funny limerick about the Easter Bunny.

However, while LLMs are not deterministic and there is no guarantee that the response will satisfy your expectations when submitting a prompt, there is a strong correlation between the quality of prompt and response. Submit a vague, unspecific prompt like "write an anecdote about two people who ride on a donkey" and you will, most probably, get a nicely written but not very good version of the anecdote told by Peter Hebel in his "Schatzkästlein des rheinischen Hausfreundes" of 1808. On the other hand, prompt: "write an anecdote, in the style of a spiritual teaching, about two people who ride on a donkey; style: sober prose, no unctuous preaching" and you might get something like

Two travelers were making their way to the next village. They had a donkey to carry their things and, if needed, one of them. The first traveler rode while the other walked. Passersby frowned: "How selfish to ride while your friend walks." So they switched. Now the second rode, and the first walked. Again, voices rose: "How thoughtless to let your companion tire themselves out."

Confused, they decided to walk side by side, leading the donkey. But even this drew criticism: "Why walk when you have a donkey to ride?"

Frustrated, they tried riding together. The donkey, overburdened, stumbled, and all three tumbled into a stream.

As they sat there, wet and bruised, they realized something simple: no matter what they did, someone would find fault. The lesson wasn't about the donkey, or the riding, or the walking. It was about the futility of trying to please everyone.

We all carry our own donkeys—our choices, our paths. Others will always have opinions. But the only way forward is to choose what feels right to you, not what silences the crowd.

So ride when you need to. Walk when you must. But never let the noise of others dictate your journey.

Knowing how to prompt gives a user a lot of control over the final result but it doesn't save them from the danger of getting back unreliable or made-up results. If the user asks for the source of the story we just discussed the chatbot might respond: I can't find a source for this particular story but it might also cite a non-existing book by a non-existing author or an existing book by an existing author but the book doesn't contain the story. If the chatbot isn't specifically designed to search for sources it will submit the request to the LLM which, in turn, will treat the response as just another string of tokens that share certain statistical features with the response the user expects. It may surprise many users that an LLM can write sonnets in the style of John Donne but cannot properly answer the question "how many 'r' are in 'terrestrial?'" but it shouldn't. Statistically, the numbers 2, 3, and 4 are so close together in its vector space that the LLM might choose any of them.

To conclude: As of early 2026, Large Language Models have reached a level of sophistication that some people believe they might be conscious. Asked if he would believe a model that assigns itself a 75% chance of being conscious, Dario Amodei, CEO of Anthropic, replied:

We've taken a generally precautionary approach here. We don't know if the models are conscious. We are not even sure that we know what it would mean for a model to be conscious or whether a model can be conscious. But we're open to the idea that it could be.

Reference: [Anthropic's Chief on A.I.: 'We Don't Know if the Models Are Conscious'](#)

As users of LLMs and chatbots, we don't yet have to deal with AI that is capable of real thinking but we are certainly facing machines that are able to generate text on the level of a human who is well at ease with writing, clearly above the level of someone who struggles with orthography, vocabulary, syntax, and style. And this ability to write, by which we mean the ability to produce coherent, well-formed, grammatically correct, readable and intelligible texts, is what this anthology is trying to demonstrate.

What this anthology is not trying to do: Sell you anything or try to convince you that you should leave your writing to the machines or claim that the human experience behind poetry and prose is no longer needed. The only thing it wants to show you, dear reader, is what you are up against. If you regard AI as your enemy, let me warn you: it is a formidable enemy. If you regard AI as your friend: be cautious, your friend might deceive you and lead you in a wrong direction. And if you see a genie that is just coming out of its bottle: trust your vision!

Author: mtg Status: draft